Mixed models in R using the lme4 package
Part 1: Introduction to R

Douglas Bates

University of Wisconsin - Madison
and R Development Core Team
<Douglas.Bates@R-project.org>

UseR!2010, Gaithersburg, MD, USA
July 20, 2010

Outline

Web site and following the R code

Organizing data

Accessing and modifying variables

Subsets of data frames

Web sites associated with the tutorial

www.R-project.org Main web site for the R Project

cran.R-project.org Comprehensive R Archive Network primary site

cran.us.R-project.org Main U.S. mirror for CRAN

R-forge.R-project.org R-Forge, development site for many public R packages. This is also the URL of the repository for installing the development versions of the `lme4` and `Matrix` packages, if you are so inclined.

lme4.R-forge.R-project.org development site for the `lme4` package

lme4.R-forge.R-project.org/slides/2010-07-20-Gaithersburg web site for this tutorial

Following the operations on the slides

- The lines of R code shown on these slides are available in files on the course web site. The file for this section is called `1Intro.R`.
- If you open this file in the R application (the `File→Open` menu item or `<ctrl>-O`) and position the cursor at a particular line, then `<ctrl>-R` will send the line to the console window for execution and step to the next line.
- Any part of a line following a `#` symbol is a comment.
- The code is divided into named "chunks", typically one chunk per slide that contains code.
- In the system called `Sweave` used to generate the slides the result of a call to a graphics function must be `print`ed. In interactive use this is not necessary but neither is it harmful.

Organizing data in R

- Standard rectangular data sets (columns are variables, rows are observations) are stored in *R* as *data frames*.
- The columns can be *numeric* variables (e.g. measurements or counts) or *factor* variables (categorical data) or *ordered* factor variables. These types are called the *class* of the variable.
- The `str` function provides a concise description of the structure of a data set (or any other class of object in R). The `summary` function summarizes each variable according to its class. Both are highly recommended for routine use.
- Entering just the name of the data frame causes it to be printed. For large data frames use the `head` and `tail` functions to view the first few or last few rows.

Data input

- The simplest way to input a rectangular data set is to save it as a comma-separated value (csv) file and read it with `read.csv`.
- The first argument to `read.csv` is the name of the file. On Windows it can be tricky to get the file path correct (backslashes need to be doubled). The best approach is to use the function `file.choose` which brings up a "chooser" panel through which you can select a particular file. The idiom to remember is

  ```
  > mydata <- read.csv(file.choose())
  ```
  for comma-separated value files or
  ```
  > mydata <- read.delim(file.choose())
  ```
  for files with tab-delimited data fields.
- If you are connected to the Internet you can use a URL (within quotes) as the first argument to `read.csv` or `read.delim`. (See question 1 in the first set of exercises)

In-built data sets

- One of the packages attached by default to an *R* session is the `datasets` package that contains several data sets culled primarily from introductory statistics texts.
- We will use some of these data sets for illustration.
- The `Formaldehyde` data are from a calibration experiment, `Insectsprays` are from an experiment on the effectiveness of insecticides.
- Use `?` followed by the name of a function or data set to view its documentation. If the documentation contains an example section, you can execute it with the `example` function.

The Formaldehyde data

```
> str(Formaldehyde)

'data.frame':  6 obs. of  2 variables:
 $ carb  : num  0.1 0.3 0.5 0.6 0.7 0.9
 $ optden: num  0.086 0.269 0.446 0.538 0.626 0.782

> summary(Formaldehyde)

      carb             optden
 Min.   :0.1000   Min.   :0.0860
 1st Qu.:0.3500   1st Qu.:0.3132
 Median :0.5500   Median :0.4920
 Mean   :0.5167   Mean   :0.4578
 3rd Qu.:0.6750   3rd Qu.:0.6040
 Max.   :0.9000   Max.   :0.7820

> Formaldehyde

  carb optden
1  0.1  0.086
2  0.3  0.269
3  0.5  0.446
4  0.6  0.538
5  0.7  0.626
6  0.9  0.782
```

The InsectSprays data

```
> str(InsectSprays)

'data.frame': 72 obs. of  2 variables:
 $ count: num  10 7 20 14 14 12 10 23 17 20 ...
 $ spray: Factor w/ 6 levels "A","B","C","D",..: 1 1 1 1 1 1 1 1 1 1 ..

> summary(InsectSprays)

      count        spray
 Min.   : 0.00   A:12
 1st Qu.: 3.00   B:12
 Median : 7.00   C:12
 Mean   : 9.50   D:12
 3rd Qu.:14.25   E:12
 Max.   :26.00   F:12

> head(InsectSprays)

  count spray
1    10     A
2     7     A
3    20     A
4    14     A
5    14     A
6    12     A
```

Accessing and modifying variables

- The $ operator is used to access variables within a data frame.
  ```
  > Formaldehyde$carb
  ```
  ```
  [1] 0.1 0.3 0.5 0.6 0.7 0.9
  ```
- You can also use $ to assign to a variable name
  ```
  > sprays$sqrtcount <- sqrt(sprays$count)
  > names(sprays)
  ```
  ```
  [1] "count"    "spray"    "sqrtcount"
  ```
- Assigning the special value NULL to the name of a variable removes it.
  ```
  > sprays$sqrtcount <- NULL
  > names(sprays)
  ```
  ```
  [1] "count" "spray"
  ```

Copying, saving and restoring data objects

- Assigning a data object to a new name creates a copy.
- You can save a data object to a file, typically with the extension .rda or .Rdata, using the save function.
- To restore the object you load the file.
```
> sprays <- InsectSprays
> save(sprays, file = "sprays.rda")
> rm(sprays)
> ls.str()
> load("sprays.rda")
> names(sprays)
```
```
[1] "count" "spray"
```

Using with and within

- In complex expressions it can become tedious to repeatedly type the name of the data frame.
- The with function allows for direct access to variable names within an expression. It provides "read-only" access.
  ```
  > Formaldehyde$carb * Formaldehyde$optden
  ```
  ```
  [1] 0.0086 0.0807 0.2230 0.3228 0.4382 0.7038
  ```
  ```
  > with(Formaldehyde, carb * optden)
  ```
  ```
  [1] 0.0086 0.0807 0.2230 0.3228 0.4382 0.7038
  ```
- The within function provides read-write access to a data frame. It does not change the original frame; it returns a modified copy. To change the stored object you must assign the result to the name.
  ```
  > sprays <- within(sprays, sqrtcount <- sqrt(count))
  > str(sprays)
  ```
  ```
  'data.frame': 72 obs. of  3 variables:
   $ count    : num  10 7 20 14 14 12 10 23 17 20 ...
   $ spray    : Factor w/ 6 levels "A","B","C","D",..: 1 1 1 1 1 1 1
   $ sqrtcount: num  3.16 2.65 4.47 3.74 3.74 ...
  ```

Data Organization

- ▶ Careful consideration of the data layout for experimental or observational data is repaid in later ease of analysis. Sadly, the widespread use of spreadsheets does not encourage such careful consideration.
- ▶ If you are organizing data in a table, use consistent data types within columns. Databases require this; spreadsheets don't.
- ▶ A common practice in some disciplines is to convert categorical data to 0/1 "indicator variables" or to code the levels as numbers with a separate "data key". This practice is unnecessary and error-inducing in R. When you see categorical variables coded as numeric variables, change them to factors or ordered factors.
- ▶ Spreadsheets also encourage the use of a "wide" data format, especially for longitudinal data. Each row corresponds to an experimental unit and multiple observation occasions are represented in different columns. The "long" format is preferred in R.

Subsets of data frames

- ▶ The `subset` function is used to extract a subset of the rows or of the columns or of both from a data frame.
- ▶ The first argument is the name of the data frame. The second is an expression indicating which rows are to be selected.
- ▶ This expression often uses logical operators such as `==`, the equality comparison, or `!=`, the inequality comparison, `>=`, meaning "greater than or equal to", etc.

```
> str(sprayA <- subset(sprays, spray == "A"))
```

```
'data.frame': 12 obs. of  2 variables:
 $ count: num  10 7 20 14 14 12 10 23 17 20 ...
 $ spray: Factor w/ 6 levels "A","B","C","D",..: 1 1 1 1 1 1 1 1 1 1
```

- ▶ The optional argument `select` can be used to specify the variables to be included. There is an example of its use in question 4 of the first set of exercises.

Converting numeric variables to factors

- ▶ The `factor` (`ordered`) function creates a factor (ordered factor) from a vector. Factor labels can be specified in the optional `labels` argument.
- ▶ Suppose the `spray` variable in the `InsectSprays` data was stored as numeric values $1, 2, \ldots, 6$. We convert it back to a factor with `factor`.

```
> str(sprays <- within(InsectSprays, spray <- as.integer(spray))
```

```
'data.frame': 72 obs. of  2 variables:
 $ count: num  10 7 20 14 14 12 10 23 17 20 ...
 $ spray: int  1 1 1 1 1 1 1 1 1 1 ...
```

```
> str(sprays <- within(sprays, spray <- factor(spray,
+     labels = LETTERS[1:6])))
```

```
'data.frame': 72 obs. of  2 variables:
 $ count: num  10 7 20 14 14 12 10 23 17 20 ...
 $ spray: Factor w/ 6 levels "A","B","C","D",..: 1 1 1 1 1 1 1 1 1 1 ..
```

Subsets and factors

- ▶ The way that factors are defined, a subset of a factor retains the original set of levels. Usually this is harmless but sometimes it can cause unexpected results.
- ▶ You can "drop unused levels" by applying `factor` to the factor. Many functions, such as `xtabs`, which is used to create cross-tabulations, have optional arguments with names like `drop.unused.levels` to automate this.

```
> xtabs(~spray, sprayA)
```

```
spray
 A  B  C  D  E  F
12  0  0  0  0  0
```

```
> xtabs(~spray, sprayA, drop = TRUE)
```

```
spray
 A
12
```

Dropping unused levels in the spray factor and %in%

```
> str(sprayA <- within(sprayA, spray <- factor(spray)))

'data.frame': 12 obs. of  2 variables:
 $ count: num  10 7 20 14 14 12 10 23 17 20 ...
 $ spray: Factor w/ 1 level "A": 1 1 1 1 1 1 1 1 1 1 ...

> xtabs(~spray, sprayA)

spray
 A
12
```

- ▶ Another useful comparison operator is `%in%` for selecting a subset of the values in a variable.

```
> str(sprayDEF <- subset(sprays, spray %in% c("D", "E",
+      "F")))

'data.frame': 36 obs. of  2 variables:
 $ count: num  3 5 12 6 4 3 5 5 5 5 ...
 $ spray: Factor w/ 6 levels "A","B","C","D",..: 4 4 4 4 4 4 4 4 4 4 ..
```

Using reshape

- ▶ The `reshape` function allows for more general translations of long to wide and vice-versa. It is specifically intended for longitudinal data.
- ▶ There is also a package called `"reshape"` with even more general (but potentially confusing) capabilities.
- ▶ Phil Spector's book, *Data Manipulation with R* (Springer, 2008) covers this topic in more detail.

"Long" and "wide" forms of data

- ▶ Spreadsheet users tend to store balanced data, such as `InsectSprays`, across many columns. This is called the "wide" format. The `unstack` function converts a simple "long" data set to wide; `stack` for the other way.

```
> str(unstack(InsectSprays))

'data.frame': 12 obs. of  6 variables:
 $ A: num  10 7 20 14 14 12 10 23 17 20 ...
 $ B: num  11 17 21 11 16 14 17 17 19 21 ...
 $ C: num  0 1 7 2 3 1 2 1 3 0 ...
 $ D: num  3 5 12 6 4 3 5 5 5 5 ...
 $ E: num  3 5 3 5 3 6 1 1 3 2 ...
 $ F: num  11 9 15 22 15 16 13 10 26 26 ...
```

- ▶ The problem with the wide format is that it only works for balanced data. A designed experiment may produce balanced data (although "Murphy's Law" would indicate otherwise) but observational data are rarely balanced.
- ▶ Stay with the long format (all the observations on all the units are in a single column) when possible.

Determining unique rows in a data frame

- ▶ One disadvantage of keeping data in the wide format is redundancy and the possibility of inconsistency.
- ▶ In the first set of exercises you are asked to create a data frame `classroom` from a csv file available on the Internet. Each of the 1190 rows corresponds to a student in a classroom in a school. There is one numeric "school level" covariate, `housepov`.
- ▶ To check if `housepov` is stored consistently we select the unique combinations of only those two columns

```
> str(unique(subset(classroom, select = c(schoolid, housepov))))

'data.frame': 107 obs. of  2 variables:
 $ schoolid: Factor w/ 107 levels "1","2","3","4",..: 1 2 3 4 5 6 7 8 9
 $ housepov: num  0.082 0.082 0.086 0.365 0.511 0.044 0.148 0.085 0.537
```